

Carter Blair

Website | carterblair@g.harvard.edu

EDUCATION

PhD, Computer Science

Harvard University

Cambridge, MA, USA

September 2025 - Present

- Supervised by Ariel Procaccia and Milind Tambe

Master of Math, Computer Science (Thesis Based)

University of Waterloo

Waterloo, Canada

September 2023 - April 2025

- Supervised by Kate Larson and Edith Law
- Thesis: [Deliberative Machines: From Reflective Dialogue to Fair Consensus with Language Models and Social Choice](#)

Bachelor of Science, Computer Science and Psychology, Minor in Philosophy

University of Victoria

Victoria, Canada

Completed August 2022

- Supervised by Nishant Mehta and George Tzanetakis
- Thesis: Online Positive Unlabelled Learning

PROFESSIONAL EXPERIENCE

Research Intern

[Vector Institute](#)

January 2024 – Present

Toronto, Canada

- Supervised by Gillian Hadfield
- Developed methods for AI alignment which rely on existing normative infrastructure and institutions.

Founder, Technical Lead

[redy.ai](#)

May 2023 – September 2024

- Conducted user research with real estate agents (n=75) to refine the problem definition.
- Developed a full-stack prototype with React, Django, and an LLM-based agent for task automation.

Junior Data Scientist

[NannyML](#)

July 2022 – May 2023

Leuven, Belgium

- Developed novel, linear time multivariate drift detection methods based on the Maximum Mean Discrepancy.
- Implemented standard drift detection methods in our open-source library, including Earth Mover's distance, Hellinger distance, Jensen-Shannon distance, and L-Infinity distance.

AWARDS AND HONORS

Canada Graduate Scholarship (CGS-D)

National Science and Engineering Research Council of Canada

2025 - 2028

\$120,000

- Awarded CGS-D (Highest ranked applicants).
- Declined the award to hold the PGS-D (valued at \$120,000) in the U.S.

Presidents Graduate Scholarship (Declined)

University of Waterloo

2025 - 2026

\$10,000

[Cooperative AI PhD Fellowship \(Scholar\)](#)

[Cooperative AI Foundation](#)

2025 - 2028

Research Dissemination Award

University of Waterloo

2024

\$500

Math Domestic Graduate Student Award

University of Waterloo

2023

\$2,000

Undergraduate Student Research Award <i>National Science and Engineering Research Council of Canada</i>	2021 \$8,000
Dean's Entrance Scholarship <i>University of Victoria</i>	2017 \$2,000

PUBLICATIONS

Conference Papers

1. **Blair, C.**, & Larson, K. (2026). Generating Fair Consensus Statements with Social Choice on Token-Level MDPs. In *Proceedings of the International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2026)*, to appear.
2. **Blair, C.**, Larson, K., & Law, E. (2025). Reflective Verbal Reward Design for Pluralistic Alignment. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence (IJCAI-25)* (Human-Centred AI track), pp. 10271–10279. doi:10.24963/ijcai.2025/1141.
3. **Blair, C.**, Armstrong, B., & Larson, K. (2024). Liquid Ensemble Selection for Continual Learning. In *Proceedings of the Canadian Conference on Artificial Intelligence (Canadian AI 2024)*.
4. **Blair, C.**, Wang, X., & Perin, C. (2024). Quantifying Emotional Responses to Immutable Data Characteristics and Designer Choices in Data Visualizations. *IEEE Transactions on Visualization and Computer Graphics (TVCG)*.

Peer-reviewed Workshop Papers, Posters, and Abstracts

1. Caiata, J., **Blair, C.**, & Larson, K. (2026). Procedural Fairness in Multi-Agent Bandits. Extended abstract in *AAMAS 2026*, to appear.
2. **Blair, C.**, Armstrong, B., Alouf-Heffetz, S., Talmon, N., & Grossi, D. (2025). Probably Approximately Consensus: On the Learning Theory of Finding Common Ground. *2nd Workshop on Social Choice and Learning Algorithms (SCaLA) at IJCAI 2025*.
3. Trivedi, R., **Blair, C.**, Chandak, N., Sarkar, A., Weltman, T., Hadfield-Menell, D., & Hadfield, G. K. (2024). Altered Environments: The Role of Normative Infrastructure in AI Alignment. *Agentic Markets Workshop at ICML 2024*. Oral presentation.
4. Sarkar, A., Muresanu, A., **Blair, C.**, Sharma, A., Trivedi, R. S., & Hadfield, G. K. (2024). Normative Modules: A Generative Agent Architecture for Learning Norms that Supports Multi-Agent Cooperation. *Foundation Models and Game Theory Workshop at Economics and Computation (EC) 2024*. Oral presentation.
5. **Blair, C.**, & Wyeth, C. (2024). Decision Theoretic Planning with Language Models. *Foundation Models and Game Theory Workshop at Economics and Computation (EC) 2024*. Poster.

PRESENTATIONS

- *Representative Preference Learning*, Google DeepMind, June 2024

TEACHING EXPERIENCE

Teaching Assistant, CS 480 - Machine Learning <i>University of Waterloo</i>	September 2024 – December 2024
Teaching Assistant, CS 137 - Programming Principles <i>University of Waterloo</i>	September 2023 – December 2023
Teaching Assistant, CS 110 - Fundamentals of Programming I <i>University of Victoria</i>	September 2021 – December 2021

REVIEWING

Reviewer <i>7th Intl. Workshop on Agents for Societal Impact at AAMAS 2026</i>	2026
Reviewer <i>IJCAI-ECAI 2026 Special Track on Human Centred AI</i>	2026
Reviewer <i>Social Choice and Learning Algorithms Workshop at IJCAI 2025</i>	2025
Reviewer <i>IJCAI 2025</i>	2025
Reviewer <i>Conference on Human Factors in Computing Systems (CHI) Late Breaking Work</i>	2024

LEADERSHIP

Waterloo Team Lead <i>Cooperative AI 2023 MeltingPot Competition</i>	September 2013 – November 2023
President <i>UVic Neuro-Tech Club</i>	September 2019 – April 2020 Victoria, BC
President <i>UVic Surf Club</i>	September 2018 – December 2019 Victoria, BC